

Zero Trust Architecture Sandbox: Implementation and Evaluation Report

Minsoo Ahn

CSE 487 T92

Student ID: 114743792

Instructor: Antonino Mione

May 23, 2026

Abstract—This report describes a hands-on implementation of a Zero Trust Architecture (ZTA) using an open-source stack on a local Kubernetes cluster. The system follows NIST SP 800-207 [1] principles with Istio as the policy enforcement point, OPA as the policy decision point, and Keycloak as the identity provider. Five security scenarios are validated: lateral movement defense, JWT forgery and tampering detection, context-based access control, JWT claim-based authorization, and device posture gating. All 18 test cases pass in end-to-end runs, and evidence is captured through automated `make test-all` logs, OPA decision logs, and a Kiali topology graph. Performance measurements show an average latency increase of 1.53 ms and a 22.3% throughput reduction under single-connection load tests. The results demonstrate that ZTA enforces per-request, context-aware authorization and workload identity controls that traditional firewalls and VPNs do not provide, while introducing a measurable but acceptable overhead for a security-focused environment.

I. INTRODUCTION

Modern breaches commonly follow a two-stage pattern: credential theft followed by lateral movement inside the network. When the stolen account is non-privileged, a privilege-escalation step—using techniques such as pass-the-hash or misconfigured `sudo`—typically precedes lateral movement, making the full kill chain three stages. Traditional perimeter controls such as firewalls and VPNs can block unauthenticated ingress but provide limited visibility and control for east-west (lateral, intra-cluster) traffic between workloads. Zero Trust Architecture changes the assumption from “trusted internal” to “verify every request” [1], [2]. This project implements a ZTA sandbox to validate that request-level identity and policy enforcement can block a set of representative attacks while remaining practical to operate on a small cluster.

II. METHODOLOGY

A. Environment and Stack

The sandbox runs on Minikube with Istio 1.28.3 for service mesh enforcement, OPA for policy evaluation, and Keycloak for JWT issuance. A simple Flask frontend and backend are deployed as separate services. Security policies are applied through Istio AuthorizationPolicy and PeerAuthentication resources, and Rego policies in OPA. The full validation suite is driven by `make test-all`, which dispatches each case through a shared runner against three long-lived client




Fig. 1. Request flow and enforcement layers. Top row: north-south path (Client → Envoy JWT verify → OPA → Frontend → Backend). Bottom row: east-west path (Frontend → mTLS SVID → SPIFFE Allowlist → Backend). OPA participates only in the north-south path; east-west enforcement relies solely on mTLS peer authentication and the SPIFFE allowlist.

workloads—a sidecar-injected client, a sidecar-free “rogue” client, and a client bound to the wrong ServiceAccount—reused via `kubectl exec` rather than spawning a fresh pod per case. This keeps the 18-case suite at roughly twenty seconds end-to-end and makes the security outcomes deterministically reproducible.

B. Policy Design

Four independent layers enforce access decisions:

- mTLS (mutual TLS) workload identity for pod-to-pod traffic.
- SPIFFE (Secure Production Identity Framework For Everyone) allowlist to restrict backend access to the frontend service account.
- JWT signature validation using Keycloak JWKS (JSON Web Key Set).
- Context-aware authorization in OPA (role, method, path, posture).

Figure 1 illustrates both the north-south (external-to-internal) and east-west enforcement flows.

C. Threat Model and Scenarios

Five scenarios represent realistic attack paths across north-south and east-west traffic. Table I maps each scenario to a STRIDE threat category, the attacker’s assumed position and credential scope, and the mitigating policy control. STRIDE—an acronym for Spoofing, Tampering, Repudiation, Information Disclosure, Denial of Service, and Elevation of Privilege—is a threat-modeling framework originally developed by Microsoft to systematically identify security threats.

TABLE I
THREAT MODEL: STRIDE MAPPING AND ATTACKER PROFILES PER SCENARIO

Sc.	STRIDE	Attacker profile (position / credentials)	Mitigating control
A	Spoofing	Internal pod / no workload identity	mTLS STRICT + SPIFFE allowlist
B	Tampering, Spoofing	External / forged or tampered JWT	JWKS signature verification
C	Elevation of Privilege	Authenticated external / unsecured role header	OPA role + method + path rule
D	Spoofing	Authenticated external / signed JWT, wrong role claim	Keycloak JWT claim + OPA
E	Elevation of Privilege	Authenticated external / valid admin JWT, risky device	OPA device posture gate

D. Scenario Selection Rationale

The scenarios were chosen to reflect common attacker behavior observed in real microservice environments and to map cleanly to NIST SP 800-207 tenets. They are representative, repeatable on a local cluster, and span both north-south and east-west traffic paths without depending on rare or exotic conditions.

- A: Lateral movement is a dominant post-compromise behavior and validates workload identity plus service-to-service authorization.
- B: Forged or tampered JWTs are realistic API abuses and test cryptographic authentication at the infrastructure layer.
- C: Context abuse (method/path) reflects over-privileged access in practice and demonstrates fine-grained authorization beyond identity.
- D: Claim-based authorization closes header spoofing and validates signed identity as the policy source of truth.
- E: Device posture models access from risky or unmanaged devices, showing that identity alone is not sufficient.

E. Authorization Policy Logic

Table II summarizes the core allow rules in the OPA Rego policy (package `istio.authz`). All rules default to deny; a request is allowed only when every condition in one rule is satisfied. The `jwt_has_role` helper calls `io.jwt.decode` to extract `payload.realm_access.roles` without re-verifying the signature, relying on Istio’s prior JWKS verification at the mesh layer. The admin rule also requires `posture_ok` to remain consistent with the device posture gate in Scenario E, preventing a privilege bypass for admin accounts when the device signal indicates risk. The header-based rules (Scenario C) are gated behind a single `demo_mode` flag in the policy; setting `demo_mode = false` removes the client-spoofable header path entirely, leaving only the cryptographically signed JWT claims (Scenarios D and E) as the production authorization source. The policy additionally exposes an audit-only `auth_reason` value (`missing / non_bearer / malformed_jwt / ok`) in OPA decision logs, so a denied

TABLE II
OPA REGO ALLOW RULES (SIMPLIFIED). [†]THE HEADER-BASED RULES ARE GATED BY `demo_mode`; WITH `demo_mode = false` ONLY THE SIGNED-JWT RULES (D, E) REMAIN ACTIVE. [‡]THE E-W RULE RESIDES IN THE REGO POLICY BUT THE EXT-AUTHZ PROVIDER SELECTOR TARGETS ONLY THE FRONTEND SERVICE; THIS RULE NEVER FIRES ON POD-TO-POD TRAFFIC, WHICH IS ENFORCED EXCLUSIVELY BY mTLS PEERAUTHENTICATION AND THE SPIFFE ALLOWLIST.

Scenario	Conditions (conjunction)	Decision
A (E-W) [‡]	<code>source.principal = frontend-sa ∧ method = GET ∧ path ⋈ /api/admin</code>	allow
A-NS-2/C-4 [†]	<code>demo_mode ∧ role header = "admin" ∧ posture_ok</code>	allow
C [†]	<code>demo_mode ∧ role header = "user" ∧ method = GET ∧ path ⋈ /api/admin</code>	allow
D (admin)	<code>jwt.realm_access.roles ∋ "admin" ∧ posture_ok</code>	allow
D (viewer)	<code>jwt.realm_access.roles ∋ "viewer" ∧ method = GET ∧ path ⋈ /api/admin</code>	allow
E	<code>jwt.realm_access.roles ∋ "admin" ∧ X-Device-Firewall ∈ {<code>⌀</code>, "enabled"}</code>	allow
Default	otherwise	deny

request can be attributed to a specific cause without that path granting access.

F. Architectural Topology Constraints

The OPA ext-authz channel operates over plaintext gRPC on port 9191 rather than mTLS. This is a deliberate architectural constraint: injecting an Envoy sidecar into the OPA pod would create a circular authorization dependency. Every Envoy sidecar in the mesh dispatches a `CheckRequest` to OPA before forwarding an application request. If OPA’s pod itself carried an injected sidecar, that sidecar’s inbound filter chain would intercept the incoming `CheckRequest` and—following the same ext-authz filter configuration—would issue a new `CheckRequest` to OPA to authorize the original authorization request. This forms an infinite loop: `Envoy → OPA → (OPA’s own Envoy intercepts) → OPA → ⋯`. OPA must therefore reside outside the sidecar mesh. A Kubernetes `NetworkPolicy` manifest (`opa-network-policy.yaml`) is provided to restrict port 9191 to intra-cluster traffic; enforcing it in production requires a `NetworkPolicy`-capable CNI such as Calico. Minikube’s default CNI (kindnet) does not enforce `NetworkPolicies`, so in this sandbox the plaintext channel is bounded only by cluster-level isolation rather than per-port access control.

III. RESULTS

A. Scenario Outcomes

All 18 test cases pass in the final validation run: Scenarios A and B produce blocked or rejected responses at the enforcement layer (mTLS/SPIFFE and JWKS verification respectively); Scenarios C, D, and E enforce context-aware authorization over authenticated requests. Detailed per-case results are shown in Table III.

B. Scenario Details

A: Lateral movement (north-south and east-west). The north-south baseline confirms unauthenticated requests are

denied, while a demo header allows access to illustrate why header-only identity is weak. For east-west movement, a rogue pod without a sidecar fails mTLS, and a pod with the wrong ServiceAccount is denied by the SPIFFE allowlist. Direct pod IP access is also blocked. These tests validate workload identity and service-to-service authorization as separate gates.

B: JWT forgery and tampering detection. A forged JWT is rejected because no valid principal is established at the mesh layer. A tampered payload with the original signature fails cryptographic verification, returning a 401. A valid Keycloak-issued JWT is accepted and passes to OPA for authorization.

C: Context-based access control. This scenario demonstrates that identity is not enough: role, method, and path must be consistent. A user role can read `/api/data` but is denied on `/api/admin` and all write operations, while an admin role can write. This uses header-based roles as a deliberate baseline for contrast with Scenario D. In the final configuration, JWT authentication is still enforced; the header-based rules illustrate why signed claims are needed even when a request is authenticated.

D: JWT claim authorization. Authorization is driven by signed claims in the JWT. An admin token can access admin endpoints, while a viewer token is limited to read-only paths. This closes header spoofing and demonstrates a cryptographically bound identity source of truth.

E: Device posture gating. Even with a valid admin JWT, access is denied when the device posture signal indicates risk. The demo uses the `X-Device-Firewall` header to model a health signal and illustrates how posture can be integrated into policy decisions.

C. Operational Evidence

OPA decision logs confirm allow and deny outcomes for each scenario, and the automated `make test-all` run provides a consolidated PASS summary. The evidence directory includes detailed logs for JWT forgery, lateral movement attempts, and device posture simulation. Figure 2 shows the Kiali workload graph confirming mTLS enforcement across active service paths.

D. Performance

Performance testing uses Fortio with 200 requests on a single persistent connection at maximum throughput (`-qps 0`). The load parameters were chosen to isolate the pure policy evaluation cost—specifically the OPA gRPC round-trip latency and JWT payload parsing—without noise from connection-establishment overhead or scheduler contention. Minikube runs inside Docker Desktop on WSL2, where CPU and memory are shared with the host OS; under concurrent load, the OS CPU scheduler must arbitrate between the WSL2 virtual machine, the Docker daemon, and other host processes. This competition for CPU time—scheduler contention—introduces timing variance that would conflate authorization latency with resource-sharing artifacts. A single-connection sequential load eliminates that noise and isolates the pure policy-evaluation

TABLE III
COMBINED TEST MATRIX: ALL SCENARIOS A–E (18 CASES). *A-NS-2 IS A DELIBERATE ALLOW CASE SHOWING THAT HEADER-ONLY IDENTITY IS EXPLOITABLE; IT IS NOT A BLOCKED ATTACK.

Scenario	Case	Request signals	Expect	Result
A	A-NS-1	GET <code>/api/admin</code> , no identity	403	403
	A-NS-2*	GET <code>/api/admin</code> , header <code>role:admin</code>	200	200
	A-EW-1	Rogue pod (no sidecar) to backend service	403/000/503	403
	A-EW-2	Wrong ServiceAccount (backend-sa) to backend	403	403
	A-EW-3	Rogue pod to backend pod IP:8080	403/000/503	503
B	B-1	Forged JWT (invalid signature) to <code>/api/admin</code>	403	403
	B-2	Tampered real JWT payload to <code>/api/admin</code>	401	401
	B-4	Valid Keycloak JWT to <code>/api/admin</code>	200	200
C	C-1	<code>role:user</code> , GET <code>/api/data</code>	200	200
	C-2	<code>role:user</code> , GET <code>/api/admin</code>	403	403
	C-3	<code>role:user</code> , POST <code>/api/write</code>	403	403
	C-4	<code>role:admin</code> , POST <code>/api/write</code>	200	200
D	D-1	Admin JWT, GET <code>/api/admin</code>	200	200
	D-2	Viewer JWT, GET <code>/api/data</code>	200	200
	D-3	Viewer JWT, GET <code>/api/admin</code>	403	403
	D-4	Viewer JWT, POST <code>/api/write</code>	403	403
E	E-1	Admin JWT + <code>X-Device-Firewall: enabled</code>	200	200
	E-2	Admin JWT + <code>X-Device-Firewall: disabled</code>	403	403

cost. The throughput figures therefore represent sequential-request capacity rather than peak concurrency. Relative to a baseline with the Istio sidecar present but no AuthorizationPolicy or PeerAuthentication applied, ZTA adds 1.53 ms to average latency (5.33 ms $\rightarrow$ 6.86 ms, +28.6%), reduces QPS (queries per second) from 187.5 to 145.7 (−22.3%), and increases p99 (99th-percentile) latency from 8.5 ms to 10.5 ms (+23.5%). Figure 3 visualizes the latency comparison.

Table IV breaks down per-scenario timing. Blocked cases measure policy-gate cost only (the request is denied before reaching the application), while allowed cases include full end-to-end handling; the two categories should be compared separately. Blocked requests average 2.32 ms versus 6.34 ms for allowed end-to-end requests—a $2.7\times$ difference—demonstrating that early policy gating terminates invalid requests before application processing, reducing per-denied-request latency by 63.4% relative to a fully handled request.




Fig. 2. Kiali Workload Graph captured during the `make for_pdf` test run, showing one representative request per scenario (A–E). The mTLS padlock icon on the frontend→backend edge confirms STRICT mutual authentication is active. Blocked paths (rogue pod, forged JWT) terminate at the ingress and do not appear as edges to the backend, illustrating that the enforcement layers are working as intended.




Fig. 3. Latency overhead introduced by ZTA policies (Fortio, 200 req, single connection at max throughput). Baseline avg 5.33 ms, ZTA avg 6.86 ms (+28.6%).

IV. DISCUSSION

A. Security Value vs Traditional Controls

The implementation demonstrates that per-request authorization at the infrastructure layer can block attack paths that traditional perimeter controls miss. Even with valid credentials, requests are denied when policy context is violated, and internal lateral movement is blocked by workload identity checks.

B. VPN vs ZTA

WireGuard [3] provides low-overhead transport security but does not enforce request-level authorization or workload identity. ZTA adds these controls with a measurable overhead, so the most practical deployment is layered: VPN for secure entry and ZTA for service policy enforcement.

TABLE IV

PER-SCENARIO LATENCY (FORTIO, 200 REQUESTS): ALLOW CASES ARE END-TO-END; BLOCK CASES MEASURE POLICY GATE ONLY. CASES A1–A4 ARE REPRESENTATIVE ALLOW PATHS; B1–B4 ARE REPRESENTATIVE BLOCK PATHS. LABELS ARE INDEPENDENT OF THE SCENARIO MATRIX IN TABLE III.

Case	Type	Request signal	Avg (ms)	p99 (ms)	Code
A1	Allow	Admin JWT, GET /api/admin	5.94	7.89	200
A2	Allow	Viewer JWT, GET /api/data	6.63	8.80	200
A3	Allow	role:user, GET /api/data	6.58	9.00	200
A4	Allow	Admin JWT + device firewall enabled	6.22	7.94	200
B1	Block	No identity, GET /api/admin	2.02	4.00	403
B2	Block	Forged JWT, GET /api/admin	2.36	5.00	403
B3	Block	Tampered JWT, GET /api/admin	2.00	4.00	401
B4	Block	Admin JWT + device firewall disabled	2.91	5.50	403

V. LIMITATIONS

Device posture is simulated through a client-provided header rather than a verified endpoint-compliance signal. The OPA ext-authz channel operates over plaintext gRPC; the architectural reason is described in Section II-F. The provided NetworkPolicy manifest is not enforced in this sandbox because Minikube’s default CNI does not support NetworkPolicy; a production deployment requires a compatible CNI (e.g., Calico) for the port-9191 restriction to take effect. Sequential-connection load tests do not capture behavior under concurrent request pressure. The environment runs on Minikube inside WSL2, where CPU and memory are shared with the host OS, so absolute throughput figures do not represent production-scale capacity.

VI. EXPERIENCE AND REFLECTION

A. Unexpected Outcomes

I expected the implementation to be harder, but the core ZTA flow was manageable once the stack was aligned. The most surprising result was performance overhead: the latency increase and QPS drop were higher than I anticipated, even in a small test.

B. Implementation Challenges

The initial setup and integration of Istio, OPA, and Keycloak was the most time-consuming part. JWT verification flow and automated end-to-end tests required careful ordering and debugging. I also spent time looking for missing scenarios and thinking about bypasses not captured by the default test set.

C. Limits of ZTA and Current Practice

This project reinforced that ZTA is a strong practical approach but not perfect. It does not eliminate risks from human

error, system misconfiguration, or novel low-level exploits, and it only blocks the scenarios it explicitly models. AI-assisted work helped accelerate policy iteration, but it did not remove the need for manual validation, and it reminded me that attackers can also use AI to discover new paths. ZTA feels like the best available option today, not a final solution.

D. Lessons Learned and Future Work

I learned that layered controls matter more than any single tool, and that the security-performance trade-off is real. As hardware improves, some overhead will be easier to absorb, but the dilemma remains. If I extend this work, I want to design a more efficient policy path and push enforcement closer to lower-level boundaries so that isolation holds even against unexpected paths.

VII. CONCLUSION

The sandbox validates that open-source tools can implement a realistic Zero Trust Architecture aligned with NIST SP 800-207 [1]. The system blocks five representative attack scenarios and provides clear audit evidence for each decision. The measured performance cost is acceptable for a security-focused environment, and the design illustrates how identity, context, and workload trust can replace implicit network trust.

REFERENCES

- [1] S. Rose, O. Borchert, S. Mitchell, and S. Connelly, “Zero Trust Architecture,” National Institute of Standards and Technology, Gaithersburg, MD, NIST SP 800-207, Aug. 2020. [Online]. Available: <https://doi.org/10.6028/NIST.SP.800-207>
- [2] R. Ward and B. Beyer, “BeyondCorp: A New Approach to Enterprise Security,” *login.*, vol. 39, no. 6, pp. 6–11, Dec. 2014. [Online]. Available: <https://research.google/pubs/pub43231/>
- [3] J. A. Donenfeld, “WireGuard: Next Generation Kernel Network Tunnel,” in *Proc. Network and Distributed System Security Symp. (NDSS)*, San Diego, CA, USA, Feb. 2017. [Online]. Available: <https://www.wireguard.com/papers/wireguard.pdf>